

Advanced cache memory optimizations

Computer Architecture

J. Daniel García Sánchez (coordinator)
David Expósito Singh
Javier García Blas

ARCOS Group
Computer Science and Engineering Department
University Carlos III of Madrid

- 1 Introduction
- 2 Advanced optimizations
- 3 Conclusion

Why do we use caching?

- To overcome the *memory wall*.
 - 1980 – 2010: Improvement in processors performance better (orders of magnitude) than memory.
 - 2005 – . . . : Situation becomes worse with emerging *multi-core* architectures.

- To reduce both data and instructions access times.
 - Make memory access time nearer to cache access time.
 - Offer the illusion of a cache size approaching to main memory size.
 - Based on the **principle of locality**.

Memory average access time

- 1 level.

$$t = t_h(L1) + m_{L1} \times t_p(L1)$$

- 2 levels.

$$t = t_h(L1) + m_{L1} \times (t_h(L2) + m_{L2} \times t_p(L2))$$

- 3 levels.

$$t = t_h(L1) + m_{L1} \times (t_h(L2) + m_{L2} \times (t_h(L3) + m_{L3} \times t_p(L3)))$$

- ...

Basic optimizations

1. Increase block size.
2. Increase cache size.
3. Increase associativity.
4. Introduce multi-level caches.
5. Give priority to read misses.
6. Avoid address translation during indexing.

Advanced optimizations

■ Metrics to be decreased:

- Hit time.
- Miss rate.
- Miss penalty.

■ Metrics to be increased:

- Cache bandwidth.

- **Observation:** All advanced optimizations aim to improve some of those metrics.

- 1 Introduction
- 2 Advanced optimizations
- 3 Conclusion

2 Advanced optimizations

- Small and simple caches
- Way prediction
- Pipelined access to cache
- Non-blocking caches
- Multi-bank caches
- Critical word first and early restart
- Write buffer merge
- Compiler optimizations
- Hardware prefetching

Small caches


- **Lookup** procedures:
 - Select a line using the **index**.
 - Read line **tag**.
 - Compare to **address tag**.

- Lookup time is **increased** as cache size grows.

- A **smaller** cache allows:
 - Simpler lookup hardware.
 - Cache can better fit into processor chip.

- **A small cache improves lookup time.**

Access time and cache size



© 2007 Elsevier, Inc. All rights reserved.

Source: Computer Architecture: A Quantitative Approach. 6 Ed
Hennessy and Patterson. Morgan Kaufmann. 2017.

Simple caches

- Cache simplification.
 - Use mapping mechanisms as **simple** as possible.
 - **Direct mapping**:
 - Allows to **parallelize** tag comparison and data transfers.

- **Observation**: Most modern processors focus more on using small caches than on simplifying them.

Intel Core i7

- L1 cache (1 per core)
 - 32 KB instructions.
 - 32 KB data.
 - Latency: 3 cycles.
 - Associative 4(i), 8(d) ways.
- L2 cache (1 per core)
 - 256 KB
 - Latency: 9 cycles.
 - Associative 8 ways.
- L3 cache (shared)
 - 8 MB
 - Latency: 39 cycles.
 - Associative 16 ways.

2 Advanced optimizations

- Small and simple caches
- **Way prediction**
- Pipelined access to cache
- Non-blocking caches
- Multi-bank caches
- Critical word first and early restart
- Write buffer merge
- Compiler optimizations
- Hardware prefetching

Way prediction

■ Problem:

- **Direct mapping** → fast but many misses.
- **Set associative mapping** → less misses but more sets (slower).

■ Way prediction

- Additional bits stored for predicting the way to be selected in the next access.
- Block prefetching and compare to single tag.
 - If there is a miss, it is compared with other tags.

2 Advanced optimizations

- Small and simple caches
- Way prediction
- **Pipelined access to cache**
- Non-blocking caches
- Multi-bank caches
- Critical word first and early restart
- Write buffer merge
- Compiler optimizations
- Hardware prefetching

Pipelined access to cache

- **Goal:** Improve cache bandwidth.
- **Solution:** Pipelined access to the cache in multiple clock cycles.
- **Effects:**
 - Clock cycle can be shortened.
 - A new access can be initiated every clock cycle.
 - Cache bandwidth is increased.
 - Latency is increased.
- Positive effect in **superscalar processors**.


2 Advanced optimizations

- Small and simple caches
- Way prediction
- Pipelined access to cache
- **Non-blocking caches**
- Multi-bank caches
- Critical word first and early restart
- Write buffer merge
- Compiler optimizations
- Hardware prefetching

Non-blocking caches

- **Problem:** Cache miss leads to a **stall** until a block is obtained.
- **Solution:** Out-of-order execution.
 - **But:** How is memory accessed while a miss is resolved?
- **Hit during miss**
 - Allow accesses with hit while waiting.
 - Reduces miss penalty.
- **Hit during several misses / Miss during miss:**
 - Allow overlapped misses.
 - Needs multi-channel memory.
 - Highly complex.

Stall rate versus maximum number of misses



© 2007 Elsevier, Inc. All rights reserved.

Source: Computer Architecture: A Quantitative Approach. 6 Ed
Hennessy and Patterson. Morgan Kaufmann. 2017.

2 Advanced optimizations

- Small and simple caches
- Way prediction
- Pipelined access to cache
- Non-blocking caches
- **Multi-bank caches**
- Critical word first and early restart
- Write buffer merge
- Compiler optimizations
- Hardware prefetching

Multi-bank caches





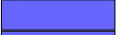











- **Goal:** Allow simultaneous accesses to different cache locations.
- **Solution:** Divide memory into independent banks.
- **Effect:** Bandwidth is increased.
- **Example:** Sun Niagara
 - L2: 4 banks.

Bandwidth

- For increasing the bandwidth, it is necessary to distribute accesses across banks.

- **Simple approach:** Sequential interleaving

- Round-robin of blocks across banks.

Block addr.	Bank 0	Block addr.	Bank 1	Block addr.	Bank 2	Block addr.	Bank 2
0		1		2		3	
4		5		6		7	
8		9		10		11	
12		13		14		15	

2 Advanced optimizations

- Small and simple caches
- Way prediction
- Pipelined access to cache
- Non-blocking caches
- Multi-bank caches
- **Critical word first and early restart**
- Write buffer merge
- Compiler optimizations
- Hardware prefetching

Critical word first and early restart

- **Observation:** Usually processors need a single word to proceed.
- **Solution:** Do not wait until the whole block from memory has been transferred.
- **Alternatives:**
 - **Critical word first:** Reorder blocks so that first word is the word needed by the processor.
 - **Early restart:** Block received without reordering.
 - As soon as the selected word is received, the processor proceeds.
- **Effects:** Depends on block size → the larger the better.

2 Advanced optimizations

- Small and simple caches
- Way prediction
- Pipelined access to cache
- Non-blocking caches
- Multi-bank caches
- Critical word first and early restart
- **Write buffer merge**
- Compiler optimizations
- Hardware prefetching

Write buffer

- A **write buffer** allows to **decrease** miss penalty.
 - When processor writes on buffer, it considers write is completed.
 - Simultaneous writes on memory are more efficient than a single write.

- **Uses:**
 - **Write-through:** On every write.
 - **Write-back:** When block is replaced.

Merges in write buffer

- If buffer contains modified blocks, addresses are checked and, if it is possible, processor performs overwrite.
- **Effects:**
 - Decrease number of **memory writes**.
 - Decrease amount of **stalls** due to full buffer.

Merges in write buffer

Write address	V		V		V		V	
100	1	M[100]	0		0		0	
108	1	M[108]	0		0		0	
116	1	M[116]	0		0		0	
124	1	M[124]	0		0		0	

Write address	V		V		V		V	
100	1	M[100]	0	M[108]	0	M[116]	0	M[124]
	1		0		0		0	
	1		0		0		0	
	1		0		0		0	

2 Advanced optimizations

- Small and simple caches
- Way prediction
- Pipelined access to cache
- Non-blocking caches
- Multi-bank caches
- Critical word first and early restart
- Write buffer merge
- **Compiler optimizations**
- Hardware prefetching

Compiler optimizations

- **Goal:** Generate code with a reduced number of instructions and data misses.

- **Instructions:**


1. Procedure reordering.
2. Align code blocks to cache line start.
3. Branch linearization.

- **Data:**

1. Array merge.
2. Loop interchange.
3. Loop merge.
4. Blocked access.

Procedure reordering

- **Goal:** Decrease **conflict misses** due to two concurrent procedures are mapped to the **same cache line**.
- **Technique:** Reorder procedures in memory.



Basic block alignment

- **Definition:** A **basic block** is a set of instructions sequentially executed (contains no branches).
- **Goal:** Decrease the **cache misses** possibility for sequential codes.
- **Technique:** Align the **first instruction** in a basic block with the **first word** in cache line.

Branch linearization

- **Goal:** Decrease cache misses due to conditional branches.
- **Technique:** If the compiler detects a branch is likely to be taken, it may invert condition and interchanges basic blocks in both alternatives.
 - Some compilers define extensions to hint the compiler.
 - **Example:** `gcc (__likely__)`.

Array merge

Parallel arrays

```
vector<int> key;  
vector<int> val;  
  
for (int i=0;i<max;++i) {  
    cout << key[i] << ", "  
        << val[i] << endl;  
}
```

Merged array

```
struct entry {  
    int key;  
    int val;  
};  
vector<entry> v;  
  
for (int i=0;i<max;++i) {  
    cout << v[i].key << ", "  
        << v[i].val << endl;  
}
```

- Decrease conflicts.
- Improve spatial locality.

Loop interchange

Striped accesses

```
for (int j=0; j<100; ++j) {  
    for (int i=0; i<5000; ++i) {  
        v[i][j] = k * v[i][j];  
    }  
}
```

Sequential accesses

```
for (int i=0; i<5000; ++i) {  
    for (int j=0; j<100; ++j) {  
        v[i][j] = k * v[i][j];  
    }  
}
```

- **Goal:** Improve spatial locality.
- Depends on the storage model defined by the programming language.
 - FOTRAN versus C.

Loop merge

Independent loops

```
for (int i=0; i<rows; ++i) {  
    for (int j=0; j<cols; ++j) {  
        a[i][j] = b[i][j] * c[i][j];  
    }  
}  
for (int i=0; i<rows; ++i) {  
    for (int j=0; j<cols; ++j) {  
        d[i][j] = a[i][j] + c[i][j];  
    }  
}
```

Merged loop

```
for (int i=0; i<rows; ++i) {  
    for (int j=0; j<cols; ++j) {  
        a[i][j] = b[i][j] * c[i][j];  
        d[i][j] = a[i][j] + c[i][j];  
    }  
}
```

- **Goal:** Improve temporal locality.
- **Beware:** It may decrease spatial locality.

Blocked access

Original product

```

for (int i=0; i<size; ++i) {
    for (int j=0; j<size; ++j) {
        r=0;
        for (int k=0; k<size; ++k) {
            r += b[i][k] * c[k][j];
        }
        a[i][j] = r;
    }
}

```

Blocked product

```

for (bj=0; bj<size; bj+=bsize) {
    for (bk=0; bk<size; bk+=bsize) {
        for (i=0; i<size; ++i) {
            for (j=bj; j<min(bj+bsize,size); ++j) {
                r=0;
                for (k=bk; k<min(bk+bsize,size); ++k) {
                    r += b[i][k] * c[k][j];
                }
                a[i][j] += r;
            }
        }
    }
}

```

■ **bsize**: Block factor

2 Advanced optimizations

- Small and simple caches
- Way prediction
- Pipelined access to cache
- Non-blocking caches
- Multi-bank caches
- Critical word first and early restart
- Write buffer merge
- Compiler optimizations
- Hardware prefetching

Instruction prefetching

- **Observation:** Instructions exhibit high spatial locality.

- **Instruction prefetching:**

- Read two consecutive blocks on miss.
 - Block causing the miss.
 - Next block.

- **Location:**

- Block causing the miss → **instruction cache**.
- Next block → **instruction buffer**.

Data prefetching

- **Example:** Pentium 4.
- **Data prefetching:** Allows to prefetch a 4KB page to L2 cache.
- Prefetching is invoked if:
 - 2 misses in L2 due to the same page.
 - Distance between misses lower than 256 bytes.

- 1 Introduction
- 2 Advanced optimizations
- 3 Conclusion

Summary (I)

■ Smaller and simpler caches

- **Improves**: Hit time.
- **Worsens**: Miss rate.
- **Complexity**: Very low.
- **Observation**: Widely used.

■ Way prediction:

- **Improves**: Hit time.
- **Complexity**: Low.
- **Observation**: Used in Pentium 4.

■ Pipelined access to cache:

- **Improves**: Hit time.
- **Worsens**: Bandwidth.
- **Complexity**: Low.
- **Observation**: Widely used.

Summary (II)

- **Non blocking access to cache:**
 - **Improves:** Bandwidth and miss penalty.
 - **Complexity:** Very high.
 - **Observation:** Widely used.
- **Multi-bank access to cache:**
 - **Improves:** Bandwidth.
 - **Complexity:** Low.
 - **Observation:** Used at L2 in Intel i7 L2 and Cortex A8.
- **Critical word first and early restart:**
 - **Improves:** Miss penalty.
 - **Complexity:** High.
 - **Observation:** Widely used.

Summary (III)

■ Write buffer merge:

- **Improves**: Miss penalty.
- **Complexity**: Low.
- **Observation**: Widely used.

■ Compiler optimizations:

- **Improves**: Miss rate.
- **Complexity**: Low for HW.
- **Observation**: Challenge is software.

■ Hardware prefetching:

- **Improves**: Miss penalty and miss rate.
- **Complexity**: Very high.
- **Observation**: More common for instructions than data accesses.

References

- **Computer Architecture. A Quantitative Approach**
5th Ed.
Hennessy and Patterson.
Sections: 2.1, 2.2.

- **Recommended exercises:**
 - 2.1, 2.2, 2.3, 2.8, 2.9, 2.10, 2.11, 2.12

Advanced cache memory optimizations

Computer Architecture

J. Daniel García Sánchez (coordinator)
David Expósito Singh
Javier García Blas

ARCOS Group
Computer Science and Engineering Department
University Carlos III of Madrid