

Virtualization and memory hierarchy


Computer Architecture

J. Daniel García Sánchez (coordinator)
David Expósito Singh
Javier García Blas

ARCOS Group
Computer Science and Engineering Department
University Carlos III of Madrid


- 1 Virtual memory
- 2 Policies
- 3 Page table
- 4 Virtual machines
- 5 VMM: Virtual Machine Monitors
- 6 Virtualization hardware support
- 7 Virtualization technologies
- 8 Conclusion

Limits of physical memory addressing



- All programs share a single addressing space.
 - **Physical address space.**
- There is no way to prevent a program to access a resource.

Overcoming the physical limit



- Programs run in a **normalized virtual addresses space**.
- **Address translation:**
 - Performed by hardware.
 - Managed by OS.
- **Supported features:**
 - Protection, Translation, Sharing.

Advantages of virtual memory (I)

■ Translation:

- Programs may have a **consistent view** of memory.
- Decreases **cost** of **multi-threaded** applications.
- Only the **working set** is needed in main memory.
- Dynamic structures only use the physical memory that they **really need** (e.g. stack).

Advantages of virtual memory (II)

■ Protection:

- Allows to **protect** a process from others.
- **Attributes** can be set at **page level**.
 - Read only, execution, ...
- **Kernel** data protected from programs.
- Improves protection against **malware**.

■ Sharing:

- **A page** can be **mapped** to several processes.
 - e.g. Memory mapped files.

Differences with cache

■ Replacement:

- **Cache**: Hardware controlled.
- **VM**: Software controlled.

■ Size:

- Cache size independent from address length.
- VM size dependent from address length.

Parameters

Parameter	L1 Cache L1	Virtual memory
Block size	16 – 128 bytes	4096 – 65,536 bytes
Hit time	1 – 3 cycles	100 – 200 cycles
Miss penalty	8 – 200 cycles	$10^6 - 10^7$ cycles
Access time	6 – 160 cycles	$8 \cdot 10^5 - 8 \cdot 10^6$ cycles
Transfer time	2 – 40 cycles	$2 \cdot 10^5 - 2 \cdot 10^6$ cycles
Miss rate	0.1% – 10%	0.00001% – 0.001%

- 1 Virtual memory
- 2 Policies
- 3 Page table
- 4 Virtual machines
- 5 VMM: Virtual Machine Monitors
- 6 Virtualization hardware support
- 7 Virtualization technologies
- 8 Conclusion

Four questions on memory hierarchy

1. Where can a block be placed in the upper level?
 - **Block placement.**

2. How is a block found in the upper level?
 - **Block identification.**

3. Which block must be replaced on miss?
 - **Block replacement.**

4. What happens on a write?
 - **Write strategy.**

Four questions on **virtual memory**

1. Where can a **page** be placed in **main memory**?
 - **Page placement.**

2. How is a **page** found in **main memory**?
 - **Page identification.**

3. Which **page** must be replaced on miss?
 - **Page replacement.**

4. What happens on a **write**?
 - **Write strategy.**

Where is a page placed in main memory?

- A page may be placed in **any page frame** in main memory.
 - Fully associative mapping.
- Managed by the operating system.
- **Goal: Minimize miss rate.**
 - Cannot do much with miss penalty.
 - Very high penalty due to slow magnetic disks.

How is a page found in main memory?

- Keep in main memory a **page table per process**.
 - **Mapping table** between **page identifier** and **page frame identifier**.
- Decreasing translation time.
 - **TLB**: *Translation Lookaside Buffer*.
 - Avoids accesses to page table in main memory.

Which page should be replaced on a miss?

- Replacement policy defined by Operating System.
 - Typically **LRU** (*Least-recently used*).

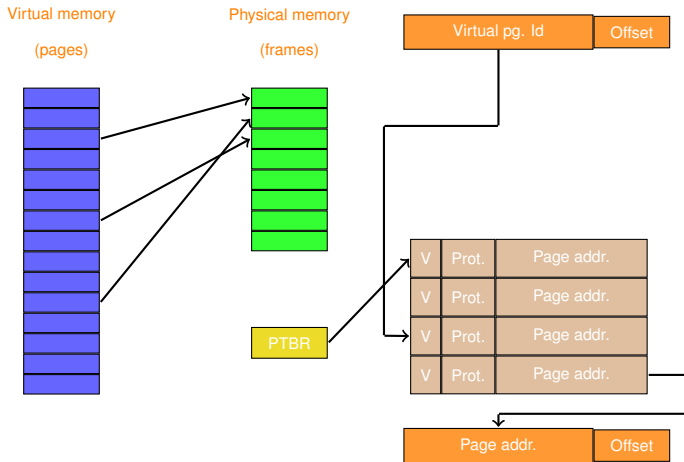
- Architecture must supply support to operating system.
 - **Use bit**: Enabled when page is accessed.
 - Really, only when TLB miss.
 - Operating system periodically zeroes this bit.
 - Records values later.
 - Allows to determine pages that have been modified within an interval.

What happens on a miss?

- Write strategy is always **write-back**.
 - No VM systems with write-through ever built.
 - **Don't be tempted!**
- Disk write costs extremely high.
 - Disk writes minimization.
 - Use **dirty bit** to annotate when a page has been modified.

- 1 Virtual memory
- 2 Policies
- 3 Page table
- 4 Virtual machines
- 5 VMM: Virtual Machine Monitors
- 6 Virtualization hardware support
- 7 Virtualization technologies
- 8 Conclusion

Page table



Page table size

- Assuming **32 bits virtual addresses**, **4 KB pages** and **4 bytes** per **table entry**:

- **Table size:**

$$\frac{2^{32}}{2^{12}} \times 2^2 B = 2^{22} B = 4MB$$

- **Alternatives:**

- Multi-level page tables.
- Inverted page tables.

- **Example:** IA-64

- Offers both alternatives to OS developer.

TLB: Translation Lookaside Buffer

■ Ideal case.

- Each access requires two memory accesses.

1. Access to page table.
2. Access to memory.

- Worse scenario in case of multi-level pages.

■ Solution:

- Use translation cache to avoid accesses to page table.
 - **Tag**: Portion of virtual address.
 - **Data**: Frame number, protection bits, validity bit, and dirty bit.

- 1 Virtual memory
- 2 Policies
- 3 Page table
- 4 Virtual machines
- 5 VMM: Virtual Machine Monitors
- 6 Virtualization hardware support
- 7 Virtualization technologies
- 8 Conclusion

Virtual machines

- Developed in late 60's.
 - Used since in *mainframe* environments.
 - Ignored in single user machines until late 90's.

- Recovered popularity due to:
 - Increasing importance of **isolation** and **security** in modern systems.
 - **Security** and **reliability** failures in operating systems.
 - **Sharing** a single computer by multiple unrelated users.
 - Dramatic increase in processor **performance**.
 - Overhead of VMMs more acceptable now.

Virtual Machine Monitor

A virtual machine is taken to be an efficient, isolated duplicate of the real machine. We explain these notions through the idea of a Virtual Machine Monitor (VMM) . . .

. . . a VMM has three essential characteristics.

- First, the VMM provides an environment for programs which is essentially identical with the original machine,
- second, programs run in this environment show at worst only minor decreases in speed;
- and last, the VMM is in complete control of system resources.

Source: Popek, G. y Goldberg, R. **Formal requirements for virtualizable third generation architectures.**

Communications of the ACM, July 1974

Virtualization

- **General definition:** Any emulation method offering a standard software interface to the physical machine.
 - Java VM? .NET?

- **System level virtual machines:** Offer a complete system environment at binary ISA level.
 - Usually assuming that VM ISA and hardware ISA are identical.
 - **Examples:**
 - IBM VM/370.
 - VMWare ESX Server.
 - Xen.

Virtual machine


- Offers to the user the illusion that they have a **complete computer**.
 - Including their own copy of the operating system.

- A computer can run **several virtual machines**.
 - May support several operating systems.
 - All operating systems sharing same hardware.

- **Terminology:**
 - **Host:** Underlying hardware platform.
 - **Guest:** Each virtual machine sharing resources.



VM y VMM: Layers



- **VMM** → Software system layer.
 - Monitor runs on hardware platform.
 - Allows execution of multiple virtual machines on single hardware platform.
 - Each virtual machine has its own operating system and applications.
 - Allows running applications without modification.

- 1 Virtual memory
- 2 Policies
- 3 Page table
- 4 Virtual machines
- 5 VMM: Virtual Machine Monitors**
- 6 Virtualization hardware support
- 7 Virtualization technologies
- 8 Conclusion

VMM

- **Virtual Machine Monitor** or **hypervisor**:
 - Software supporting virtual machines.
- VMM determines mapping between: **virtual resources** and **physical resources**.
- **Alternatives for physical resource sharing**:
 - Time sharing.
 - Partitioning.
 - Software emulation.
- A VMM is **smaller** than a traditional OS.

Overhead of VMM

- Depends on workloads.
- User level **processor bound** programs:
 - **Example**: SPEC.
 - **Overhead**: 0.
 - Invocations to OS are rare.
- **I/O intensive** programs → **OS intensive**.
 - Many system calls → Privileged instructions.
 - May lead to **a lot of virtualization overhead**.
- **I/O intensive** and **I/O bound** programs.
 - Low processor utilization.
 - **Virtualization** may be **hidden**.
 - **Low virtualization overhead**.

Other uses (in addition to protection)



■ Software management.

- VM offers an **abstraction** allowing to run a **complete software stack**.
 - Old operating Systems (DOS?).
- **Combined** deployment:
 - Stable OS, legacy OS, and next OS version.


■ Hardware management.

- VM allows to run **separate software stacks** but on top of a single hardware platform.
 - Servers consolidation.
 - Independence *rightarrow* Higher reliability.
- **Migrating** VMs in execution.
 - Load balancing.
 - Hardware evacuation due to failures.


Uses: isolation



Uses: consolidation



Uses: migration



VMM requirements (I)

- A **VMM**:
 - Offers a **software interface** to guest software.
 - **Isolates** a guest state from the rest.
 - **Protects** itself from guests.

- **Guest software** should behave as if there was no VMM, except for:
 - Performance dependent behavior.
 - Limitations of fixed resources when shared among multiple VMMs.

VMM requirements (II)

- Guest software must not be able to modify directly real resources allocation.
- VMM must control everything, even if used by guests.
 - Access to privileged state, address translation, I/O, exceptions, interruptions, ...
- VMM must run in a more privileged mode than guests.
 - Execution of privileged instructions by VMM.
- Requirements of VMM (equivalent to requirements for virtual memory).
 - A minimum of two processor modes.
 - Subset of privileged instructions, only in privileged mode.
 - Trap if executed in user mode.

- 1 Virtual memory
- 2 Policies
- 3 Page table
- 4 Virtual machines
- 5 VMM: Virtual Machine Monitors
- 6 Virtualization hardware support
- 7 Virtualization technologies
- 8 Conclusion

ISA support

- If VM considered in ISA design, it is easy to reduce instructions that VMM must execute and emulation time.
 - **But**, Most desktop ISAs designed before VMs.
- VMM must ensure that guests only interact with virtual resources.
 - Guest OS running in user mode.
 - Trying to access hardware leads to a trap.
- If ISA is not VM-aware, VMM must intercept problematic instructions.
 - Introduction of virtual resource.

Impact on virtual memory

- Each guest manages virtual memory.
 - Virtualizing virtual memory?
- VMM distinguishes between **real memory** and **physical memory**.
 - **Real Memory**: Intermediate level between **virtual memory** and **physical memory**.
 - **Guest**: Mapping between **virtual memory** and **real memory**.
 - **VMM**: Mapping between **real memory** and **physical memory**.
- To decrease indirection level, VMM keeps a **shadow page table**.
 - Mapping between **virtual memory** and **physical memory**.
 - VMM must capture changes in **page table** and **pointer to page table**.

ISA support for virtual memory virtualization

- IBM 370 includes additional indirection level managed by VMM.
 - Eliminates the need of shadow the page table.

- TLB virtualization.
 - VMM manages TLB and keeps copies of TLB in each guest.
 - TLB accesses generate a trap.
 - TLB with process identifiers simplifies management.
 - Allow entries from multiple VMs over the VMM at the same time.

Input/Output impact

- **Most complex** part in virtualization.
 - Increasing number of I/O devices.
 - Increasing diversity of I/O devices.
 - Sharing devices among VMs.
 - Support for an increasing variety of drivers.
- **General part** of driver remains on the guest side.
 - **Specific part** in VMM.
- **Device dependent** method.
 - **Disks**: Partitioned by VMM for creating virtual disks.
 - **Network interfaces**: Multiplexed over time.
 - VMM manages **virtual network addresses**.



- 1 Virtual memory
- 2 Policies
- 3 Page table
- 4 Virtual machines
- 5 VMM: Virtual Machine Monitors
- 6 Virtualization hardware support
- 7 Virtualization technologies
- 8 Conclusion

7 Virtualization technologies

- Impure virtualization
- ISA Technologies

Impure virtualization

- Solution for **non-virtualizable** architectures and for decreasing **performance problems**:
- **Approaches**:
 - **Paravirtualization**: Port guest OS code to a modified ISA.
 - Development effort.
 - Need to adapt code for every OS.
 - Source code must be available.
 - **Binary translation**: Replace non-virtualizable instructions by emulation code or VMM calls.
 - Does not require source code.
 - Some emulations are possible in user space.

Example: XEN


- **Xen**: VMM *open-source* for x86.
- **Strategy: Paravirtualization.**
 - Small modifications into the OS to simplify virtualization.
- **Examples of paravirtualization:**
 - Avoid TLB **flush** when the VMM is invoked.
 - Xen mapped to upper 64MB in each VM.
 - Allow guest to **allocate** pages.
 - Check if protection restrictions are not violated.
 - **Protection** between programs and guest → Use **protection levels** from x86:
 - Xen (0), Guest (1), Programs (3).

Changes in Xen

- Changes needed in Linux → around 3,000 lines of code.
 - 1% x86 specific code.

Operating System	Runs as <i>host</i>	Runs as <i>guest</i>
Linux 2.4	Yes	Yes
Linux 2.6	Yes	Yes
NetBSD 2.0	No	Yes
NetBSD 3.0	Yes	Yes
Plan 9	No	Yes
FreeBSD 5	No	Yes

Performance in Xen




7 Virtualization technologies

- Impure virtualization
- ISA Technologies

Intel Virtualization Technology

■ Adds new instructions:


- VMXON
- VMXOFF
- VMLAUNCH
- VMRESUME
- ...



AMD Secure Virtual Machine

■ Adds new instructions:

- VMRUN/VMLOAD.
- VMCALL/VMSAVE.
- ...



Operation modes

■ VMX root:

- Fully privileged.
- Designed to be used with VMM.

■ VMX non-root:

- Non privileged.
- Designed to be used by guest software.

Entering and exiting virtual machines

■ VM Entry:

- Transition from VMM to host.
- Enters non-root mode.
- Loads guest state.
- **VMLAUNCH** instruction used for initial entry.
- **VMRESUME** instruction used for subsequent entries.

■ VM Exit:

- **VMEXIT** instruction used to enter VMM mode.
- Enters in root mode.
- Saves guest state.
- Loads VMM state.
- There are instructions and events that cause a **VMEXIT**.

Benefits from VT technology

- Decreases OS dependency.
 - Removes needs for binary translation.
 - Facilitates the support for old operating systems.

- Improves robustness.
 - Removes need for complex techniques.
 - Smaller and simpler VMMs.

- Improves performance.
 - Less transitions to VMM.

- 1 Virtual memory
- 2 Policies
- 3 Page table
- 4 Virtual machines
- 5 VMM: Virtual Machine Monitors
- 6 Virtualization hardware support
- 7 Virtualization technologies
- 8 Conclusion**

Summary

- Virtual memory offers a mechanism for **translation**, facilitating **protection** and **sharing**.
- Virtual memory policies:
 - Placement: Fully associative.
 - Identification: Page Table.
 - Replacement: Usually LRU with TLB support.
 - Writing: Always write-back.
- Virtual machines: isolation, security, reliability, and sharing.
- Uses of VMM: protection, management sw/hw (isolation, consolidation, migration).
- Technologies: Impure virtualization and solutions in ISA.

References

- **Computer Architecture. A Quantitative Approach**
5th Ed.
Hennessy and Patterson.
Sections: B.4, 2.4.

- **Recommended exercises:**
 - B.12, B.13, B.14, 2.20, 2.21, 2.22, 2.23

Virtualization and memory hierarchy

Computer Architecture

J. Daniel García Sánchez (coordinator)
David Expósito Singh
Javier García Blas

ARCOS Group
Computer Science and Engineering Department
University Carlos III of Madrid